

# PS4\_junhoc

February 5, 2019

## 0.1 MACS30150 PS4

### 0.1.1 Dr. Richard Evans

Submitted by Junho Choi

### 0.1.2 Problem 1

**Problem 1-(a)** In the below few code chunks, I import the data from *incomes.txt* and first try to provide a few rows of the dataset. Also, it is available to us that the minimum and maximum of the dataset are approximately 49278.80 and 135865.03 dollars, respectively.

```
In [1]: import numpy as np
        from matplotlib import pyplot as plt
        import scipy.stats as sts
        import pandas as pd

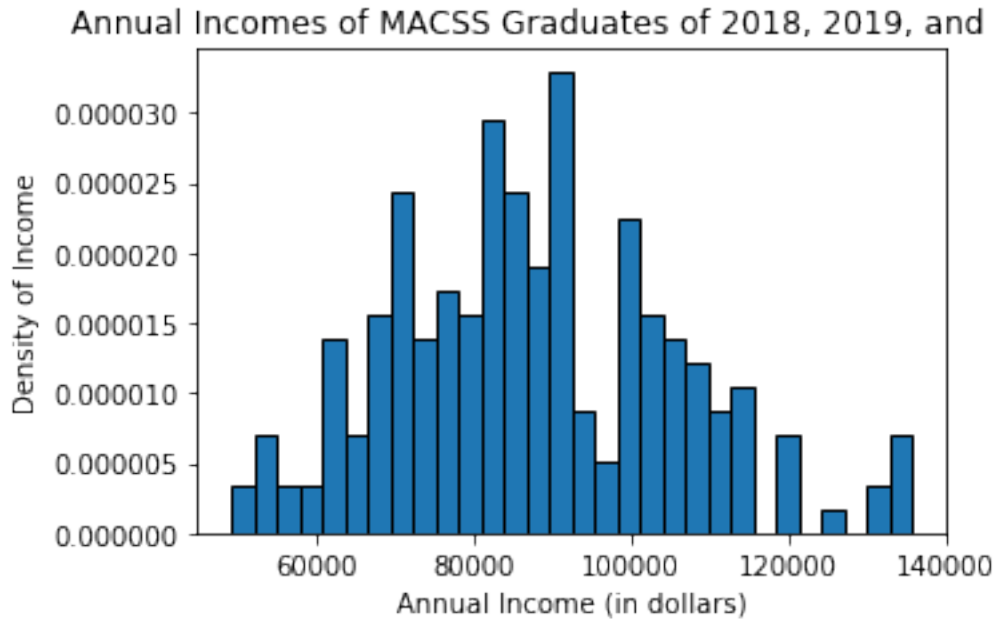
In [2]: ### Loading in the data
        y_dat = np.loadtxt('incomes.txt')
        print(y_dat[0:5]) ## checking a few elements
        print("Minimum is", '%.2f' % y_dat.min())
        print("Maximum is", '%.2f' % y_dat.max())

[ 51253.49715631 100630.32024137  83009.27613739  82882.10654304
  77338.29483892]
Minimum is 49278.80
Maximum is 135865.03
```

In the below code chunk, I plot of the histogram of annual incomes of MACSS graduates of 2018, 2019, and 2020 as requested.

```
In [3]: num_bins = 30
        count, bins, ignored = plt.hist(y_dat,
                                         num_bins, normed=True,
                                         edgecolor='k')

        plt.title('Annual Incomes of MACSS Graduates of 2018, 2019, and 2020', fontsize=12)
        plt.xlabel(r'Annual Income (in dollars)')
        plt.ylabel(r'Density of Income')
        plt.subplots_adjust(bottom=.25, left=.25)
        plt.show()
```



**Problem 1-(b)** Firstly, I prepare function *lognormalpdf* for returning the PDF density value given some point  $x$  on the distribution,  $\mu$  and  $\sigma$  where we assume that  $x \sim LN(\mu = \mu, \sigma = \sigma)$ .

```
In [4]: def lognormalpdf(x, mu, sig):

    err_msg = "standard dev. must be positive"
    assert sig > 0, err_msg

    err_msg2 = "input must be positive"
    if type(x) == np.array or type(x) == np.ndarray:
        assert x.all() >= 0, err_msg2
    elif type(x) == int or type(x) == float:
        assert x >= 0, err_msg2

    fricpart = 1 / (x * sig * ((2 * np.pi)**0.5))
    expopart = (-1) * ((np.log(x)-mu)**2) / (2*(sig**2))
    var = fricpart * (np.e ** expopart)


    return var
```

In the below code chunk, I first create a vector of equally-spaced  $x$  values for PDF graph construction -- this vector is called *onebyone* below. Notice that I plot from 0.1 to 150000.1 as 0.0 is actually not included in the support of log-normal distributions. I then plug this vector into the *lognormalpdf* function that I have created above with  $\mu = 11.0$  and  $\sigma = 0.5$ . This yields the vector of densities for log-normal distributions for the aforementioned  $x$  values.

```
In [5]: onebyone = np.linspace(0.1, 150000.1, 1500000)
        log_byone = lognormalpdf(onebyone, 11.0, 0.5)
```

Below is the resulting PDF distribution.

```
In [6]: plt.plot(onebyone, log_byone)
        plt.title("Log-normal distribution, PDF")
        plt.ylabel(r'Probability density')
        plt.xlabel(r'x')
        plt.subplots_adjust(bottom=.25, left=.25)
        plt.show()
```




Notice that the above distribution, when plotted together with the histogram, seems to be a bad fit for representing the data. Therefore, I try out another log-normal distribution with  $\mu = 11.35$  and  $\sigma = 0.2$  -- this certainly seems to be a better representation than before.

```
In [13]: ## histogram
        num_bins = 30
        count, bins, ignored = plt.hist(y_dat,
                                         num_bins, normed=True,
                                         edgecolor='k')

        log_byone = lognormalpdf(onebyone, 11.0, 0.5)
        plt.plot(onebyone, log_byone, label='mu=11.0, sig=0.5')

        log_byone2 = lognormalpdf(onebyone, 11.35, 0.2)
        plt.plot(onebyone, log_byone2, label='mu=11.35, sig=0.2')
        plt.title("Plotting of histogram and PDFs")
```

```
plt.ylabel(r'Density')
plt.xlabel(r'Annual income (in dollars)')
plt.legend()
plt.show()
```



In the below code chunk, I define the function `loglike_calc` that takes in function `fn`, `x` values (in numpy vector format) `xvals`, mean variable `mu`, and standard deviation variable `sig`.

```
In [14]: def loglike_calc(fn, xvals, mu, sig):

    msg = "please let the input be in numpy array format"
    assert type(xvals) == np.ndarray or type(xvals) == np.array, msg

    likeli_vec = fn(xvals, mu, sig)
    loglike_vec = np.log(likeli_vec)
    loglike = np.sum(loglike_vec)

    return loglike
```

Given  $\mu = 11.0$  and  $\sigma = 0.5$ , the log likelihood is found to be approximately  $-2385.857$ . On the other hand, that with  $\mu = 11.35$  and  $\sigma = 0.5$  is found to be approximately  $-2242.253$ , which certainly seems to perform better than the given set of parameters.

```
In [15]: initial_one = loglike_calc(lognormalpdf, y_dat, 11.0, 0.5)
        better_one = loglike_calc(lognormalpdf, y_dat, 11.35, 0.2)
```

```

print("With mu=11.00 and sigma=0.5, the log likelihood is", '%.3f' % initial_one)
print("With mu=11.35 and sigma=0.2, the log likelihood is", '%.3f' % better_one)

```

With mu=11.00 and sigma=0.5, the log likelihood is -2385.857

With mu=11.35 and sigma=0.2, the log likelihood is -2242.253

**Problem 1-(c)** In the below code chunk, I have created the function *crit\_lognormal* for returning the negative of log likelihood given some parameters as well as x values. This will serve as the criterion function for determining the values of  $\mu$  and  $\sigma$  that maximize log likelihood (i.e. that are maximum-likelihood estimates).

```

In [19]: def crit_lognormal(params, xvals):

    mu, sigma = params
    loglikeli = loglike_calc(lognormalpdf, xvals, mu, sigma)
    negloglikeli = -loglikeli

    return negloglikeli

```

Yet since I have not put in the x values that I will be using for this exercise, I have defined a lambda function *crit\_lognormal\_with\_x* in the code chunk below that puts in the provided income data. The maximum-likelihood estimates for  $\mu$  and  $\sigma$  are found to be  $\mu_{MLE} \approx 11.359$  and  $\sigma_{MLE} \approx 0.208$ .

```

In [20]: import scipy.optimize as opt

crit_lognormal_with_x = lambda params: crit_lognormal(params, y_dat)

mu_init = 11.35 # mu_2
sig_init = 0.2 # sig_2
params_init = np.array([mu_init, sig_init])
mle_args = y_dat
results_uncstr = opt.minimize(crit_lognormal_with_x, params_init)
mu_MLE, sig_MLE = results_uncstr.x
print('mu_MLE=', mu_MLE, ' sig_MLE=', sig_MLE)

mu_MLE= 11.3590230068 sig_MLE= 0.20817732092

```

We can certainly check below that log likelihood for MLE approach certainly is higher than the two previous ones that we had examined.

```

In [21]: initial_one = loglike_calc(lognormalpdf, y_dat, 11.0, 0.5)
better_one = loglike_calc(lognormalpdf, y_dat, 11.35, 0.2)
MLE_one = loglike_calc(lognormalpdf, y_dat, mu_MLE, sig_MLE)

print("With mu=11.00 and sigma=0.5, the log likelihood is", '%.3f' % initial_one)
print("With mu=11.35 and sigma=0.2, the log likelihood is", '%.3f' % better_one)
print("MLE log likelihood is", '%.3f' % MLE_one)

```

With  $\mu=11.00$  and  $\sigma=0.5$ , the log likelihood is -2385.857  
With  $\mu=11.35$  and  $\sigma=0.2$ , the log likelihood is -2242.253  
MLE log likelihood is -2241.719

Just in case, let us show the well-behavedness of the maximum-likelihood estimation for this problem as well. The most ideal shape of the 3-dimensional graph would be that of a paraboloid, but it seems that -- at least in the neighborhood of the maximum likelihood estimates that we had found -- the optimization depends a lot on finding the right value of  $\sigma$  and not so much with that for  $\mu$ .

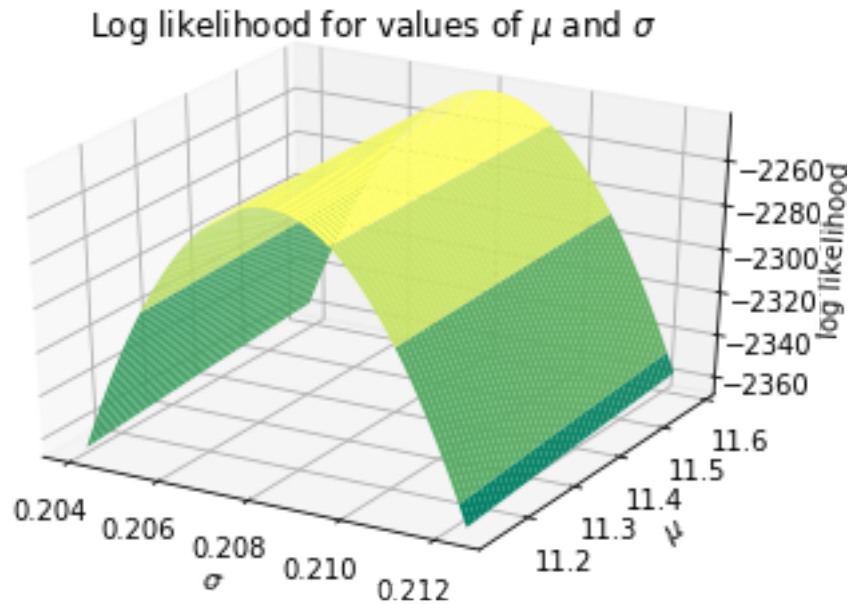
```
In [22]: import matplotlib
         from mpl_toolkits.mplot3d import Axes3D
         from matplotlib import cm
         cmap1 = matplotlib.cm.get_cmap('summer')

         # mu_vals = np.linspace(300, 635, 50)
         # sig_vals = np.linspace(30, 210, 50)
         mu_vals = np.linspace(0.98*mu_MLE, 1.02*mu_MLE, 50)
         sig_vals = np.linspace(0.98*sig_MLE, 1.02*sig_MLE, 50)
         lnlik_vals = np.zeros((50, 50))
         for mu_ind in range(50):
             for sig_ind in range(50):
                 lnlik_vals[mu_ind, sig_ind] = \
                     loglike_calc(lognormalpdf, y_dat,
                                   mu_vals[mu_ind], sig_vals[sig_ind])

         mu_mesh, sig_mesh = np.meshgrid(mu_vals, sig_vals)

         fig = plt.figure()
         ax = fig.gca(projection='3d')
         ax.plot_surface(sig_mesh, mu_mesh, lnlik_vals, rstride=8,
                          cstride=1, cmap=cmap1)
         ax.set_title(r'Log likelihood for values of  $\mu$  and  $\sigma$ ')
         ax.set_xlabel(r' $\sigma$ ')
         ax.set_ylabel(r' $\mu$ ')
         ax.set_zlabel(r'log likelihood')

Out[22]: Text(0.5,0,'log likelihood')
```



The variance-covariance matrix is also presented below.

```
In [23]: vcv_mle = results_uncstr.hess_inv
         stderr_mu_mle = np.sqrt(vcv_mle[0,0])
         stderr_sig_mle = np.sqrt(vcv_mle[1,1])
         print('Variance-Covariance Matrix of MLE:')
         print(vcv_mle)
         print()
         print('Standard error for mu estimate = ', stderr_mu_mle)
         print('Standard error for sigma estimate = ', stderr_sig_mle)
```

```
Variance-Covariance Matrix of MLE:
[[ 7.31778409e-04  3.43486900e-05]
 [ 3.43486900e-05  1.70453877e-05]]
```

```
Standard error for mu estimate = 0.0270514030802
Standard error for sigma estimate = 0.00412860602646
```

**Problem 1-(d)** The  $p$ -value for this likelihood ratio test yields 0, which means that we can reject the null hypothesis that the income data came from a log-normal distribution with  $\mu = 11.0$  and  $\sigma = 0.5$ .

```
In [24]: mu_in_b, sig_in_b = 11.0, 0.5
         log_lik_h0 = loglike_calc(lognormalpdf, y_dat, mu_in_b, sig_in_b)
         print('hypothesis value log likelihood:', log_lik_h0)
         print()
         log_lik_mle = loglike_calc(lognormalpdf, y_dat, mu_MLE, sig_MLE)
```

```

print('MLE log likelihood:', log_lik_mle)
print()
LR_val = 2 * (log_lik_mle - log_lik_h0)
print('likelihood ratio value:', LR_val)
print()
pval_h0 = 1.0 - sts.chi2.cdf(LR_val, 2)
print('chi squared of H0 with 2 degrees of freedom p-value:', pval_h0)

```

hypothesis value log likelihood: -2385.85699781

MLE log likelihood: -2241.71930136

likelihood ratio value: 288.275392902

chi squared of H0 with 2 degrees of freedom p-value: 0.0

**Problem 1-(e)** In the below code chunks, I have presented two ways -- numerical integration via the trapezoid method and directly calculating the CDF of log-normal distribution -- of calculating the percentages for earning more than 100,000 dollars and less than 75,000 dollars. The percentages are found to be approximately 22.99% for the former, and 26.02% for the latter.

In [25]: `from scipy.special import erf`

```

def lognormalcdf(x, mu, sig):
    err_msg = "standard dev. must be positive"
    assert sig > 0, err_msg

    err_msg = "input must be positive"
    if type(x) == np.array or type(x) == np.ndarray:
        assert x.all() > 0, err_msg2
    elif type(x) == int or type(x) == float:
        assert x > 0, err_msg2

    inerf = (np.log(x) - mu) / (sig * (2 ** 0.5))
    rtnval = 0.5 + 0.5 * erf(inerf)

    return rtnval

```

In [26]: `poor = lognormalcdf(75000, mu_MLE, sig_MLE)`  
`rich = 1 - lognormalcdf(100000, mu_MLE, sig_MLE)`

In [27]: `print("Probability of earning less than $75000 is", "%.2f" % (100*poor) + "%")`  
`print("Probability of earning more than $100000 is", "%.2f" % (100*rich) + "%")`

Probability of earning less than \$75000 is 26.02%  
 Probability of earning more than \$100000 is 22.99%



```

In [28]: MLE_lognorm = lambda x: lognormalpdf(x, mu_MLE, sig_MLE)

In [29]: def trapintegr(func, a, b, N):
    bin_cuts = np.linspace(a, b, N + 1)
    binsize = (b-a)/N

    bin_fvals = func(bin_cuts)
    bin_vals = (bin_fvals[1:] + bin_fvals[0:N]) / 2 * binsize

    return bin_vals.sum()

In [30]: poor2 = trapintegr(MLE_lognorm, 0.1, 75000, 75000)
    rich2 = 1 - trapintegr(MLE_lognorm, 0.1, 100000, 75000)

In [31]: print("Probability of earning less than $75000 is", "%.2f" % (100*poor2) + "%")
    print("Probability of earning more than $100000 is", "%.2f" % (100*rich2) + "%")

Probability of earning less than $75000 is 26.02%
Probability of earning more than $100000 is 22.99%

```

### 0.1.3 Problem 2

**Problem 2-(a)** To start, let us load the data and check out how the data is structured.

```

In [32]: import pandas as pd
    from sklearn.linear_model import LinearRegression as lin

In [33]: ### Loading in the data
    sick_dat = pd.read_csv('sick.txt')
    sick_dat[0:5] ## checking a few rows

Out[33]:
   sick   age  children  avgtemp_winter
0  1.67  57.47     3.04             54.10
1  0.71  26.77     1.20             36.54
2  1.39  41.85     2.31             32.38
3  1.37  51.27     2.46             52.94
4  1.45  44.22     2.72             45.90

```

To provide ourselves with a starting point for finding the maximum-likelihood estimates, let us use OLS to find the estimates for parameters.

```

In [34]: ### trying the OLS first; first prepping the y and X
    y_vals = sick_dat['sick']
    x_mat = sick_dat[['age', 'children', 'avgtemp_winter']]

    y_vals = np.array(y_vals)
    x_mat = np.array(x_mat)

```

It seems that the OLS regression estimates are as follows:  $\beta_0 \approx 0.252$ ,  $\beta_1 \approx 0.013$ ,  $\beta_2 \approx 0.400$ , and  $\beta_3 \approx -0.010$ .

```
In [35]: reg = lin().fit(x_mat, y_vals)
         base_coef = reg.coef_
         base_intercept = reg.intercept_
         print("Regression coefficients:", reg.coef_)
         print("Regression intercept:", base_intercept)

Regression coefficients: [ 0.01293366  0.40049939 -0.00999174]
Regression intercept: 0.251641374723
```

Let us now use the below function *eps\_creation* to find a vector (numpy array) of error terms (epsilons) and find the standard deviation for it. It seems that the error terms have a standard deviation of approximately 0.003.

```
In [36]: def eps_creation(y, x1, x2, x3, b0, b1, b2, b3):

         epsilons = y - b0 - b1*x1 - b2*x2 - b3*x3

         return epsilons

In [37]: epsilons = eps_creation(y_vals, x_mat[:, 0], x_mat[:, 1], x_mat[:, 2],
                                base_intercept, base_coef[0], base_coef[1], base_coef[2])

         print("Standard deviation of the error terms is", '%.5f' % epsilons.std())

Standard deviation of the error terms is 0.00302
```

Now, let us use the below function *normalpdf* to draw the normal distribution which will be presented along with the histogram for the error terms.

```
In [38]: def normalpdf(x, mu, sig):

         sig = abs(sig)

         fracpart = 1 / ((2 * np.pi * (sig ** 2)) ** 0.5)
         expopart = (-1) * ((x - mu)**2)/(2*(sig**2))
         var = fracpart * (np.e ** expopart)

         return var


In [40]: ## histogram
         num_bins = 30
         count, bins, ignored = plt.hist(epsilons,
                                         num_bins, normed=True,
                                         edgecolor='k')

         ## normal distribution
         lb, ub= np.min(epsilons), np.max(epsilons)
```

```

xspace = np.linspace(lb, ub, 1001)
norm_byone = normalpdf(xspace, 0, epsilons.std())
plt.plot(xspace, norm_byone, label='mu=0, sigma=0.003')
plt.title("Histogram of the error terms and normal distribution")
plt.ylabel(r'Density')
plt.xlabel(r'Errors')
plt.subplots_adjust(bottom=.25, left=.25)
plt.legend()
plt.show()

```



To create the log-likelihood for the normal distribution (as the error terms are assumed to be normally distributed), I use the function `ll_normal`. I admit that this function would not be suitable for those trying to use more or less than 3 regressors; in the future, I will fix this. But for now, this would be sufficient.

```

In [41]: def ll_normal(y, x1, x2, x3, b0, b1, b2, b3, sigma):

    epsilons = eps_creation(y, x1, x2, x3, b0, b1, b2, b3)
    normal_densities = normalpdf(epsilons, 0, sigma)
    ln_densities = np.log(normal_densities)
    log_ll = np.sum(ln_densities)

    return log_ll

```

I also provide the criterion function that returns the negative of the log likelihood given the parameters and arguments.

```

In [42]: def crit_normal(params, *args):

```

# PS4\_junhoc-Copy1

February 5, 2019

I also provide the criterion function that returns the negative of the log likelihood given the parameters and arguments.

```
In [42]: def crit_normal(params, *args):

        b0, b1, b2, b3, sigma = params
        y, x1, x2, x3 = args

        return -ll_normal(y, x1, x2, x3, b0, b1, b2, b3, sigma)
```

I use the initialization  $(0.2, 0.01, 0.4, -0.01, 0.05)$  which are close to the OLS parameters. The below process, for some reason, has produced programming errors; however, the estimated (using MLE approach) coefficients are exactly the same as that of the OLS, which are  $\beta_0 \approx 0.252, \beta_1 \approx 0.013, \beta_2 \approx 0.400$ , and  $\beta_3 \approx -0.010$ . The value of the log likelihood is approximately 876.87.

```
In [43]: ## base_coef and base_intercept from OLS estimation
        arguments = (y_vals, x_mat[:, 0], x_mat[:, 1], x_mat[:, 2])
        params_init2 = (0.2, 0.01, 0.4, -0.01, 0.05)

        results_uncstr2 = opt.minimize(crit_normal, params_init2, arguments)
        print(results_uncstr2.x)
        print(-results_uncstr2.fun)
```

```
[ 0.25164623  0.01293336  0.40050207 -0.00999168  0.00301767]
876.8650464975143
```

```
C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:5: RuntimeWarning: divide by zero
    """
```

```
C:\ProgramData\Anaconda3\lib\site-packages\scipy\optimize\optimize.py:643: RuntimeWarning: invalid value encountered in divide
    grad[k] = (f*((xk + d,) + args)) - f0) / d[k]
```

```
C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:5: RuntimeWarning: divide by zero
    """
```

```
C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:5: RuntimeWarning: divide by zero
    """
```

```
C:\ProgramData\Anaconda3\lib\site-packages\scipy\optimize\optimize.py:643: RuntimeWarning: invalid value encountered in divide
    grad[k] = (f*((xk + d,) + args)) - f0) / d[k]
```

```
C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:5: RuntimeWarning: divide by zero
```

"""

Below code presents the variance-covariance matrix.

```
In [44]: vcv_mle2 = results_uncstr2.hess_inv
         print('Variance-Covariance Matrix of MLE:')
         print(vcv_mle2)
```

Variance-Covariance Matrix of MLE:

```
[[ 6.83931654e-07  8.08438267e-09 -1.50571088e-07 -1.65440760e-08
   1.84103852e-09]
 [ 8.08438267e-09  3.95820755e-09 -3.52328259e-08 -2.49904455e-09
 -1.36777360e-10]
 [-1.50571088e-07 -3.52328259e-08  3.58034576e-07  2.23578038e-08
 -6.98637598e-11]
 [-1.65440760e-08 -2.49904455e-09  2.23578038e-08  1.83866092e-09
  9.79430515e-11]
 [ 1.84103852e-09 -1.36777360e-10 -6.98637598e-11  9.79430515e-11
  2.19890756e-08]]
```

**Problem 2-(b)** Below presents the results of the likelihood ratio test; because the  $p$ -value is 0.0, we can reject the null hypothesis that the regressors have no effect on number of sick days. Moreover, such a likelihood is, once again, 0%.

```
In [46]: hyp_b0, hyp_b1, hyp_b2, hyp_b3, hyp_sigma = 1, 0, 0, 0, 0.1
         mle_b0, mle_b1, mle_b2, mle_b3, mle_sigma = results_uncstr2.x

         log_lik_h0 = ll_normal(y_vals, x_mat[:, 0], x_mat[:, 1], x_mat[:, 2],
                                hyp_b0, hyp_b1, hyp_b2, hyp_b3, hyp_sigma)
         print('hypothesis value log likelihood:', log_lik_h0)
         print()
         log_lik_mle = ll_normal(y_vals, x_mat[:, 0], x_mat[:, 1], x_mat[:, 2],
                                mle_b0, mle_b1, mle_b2, mle_b3, mle_sigma)
         print('MLE log likelihood:', log_lik_mle)
         print()
         LR_val = 2 * (log_lik_mle - log_lik_h0)
         print('likelihood ratio value:', LR_val)
         print()
         pval_h0 = 1.0 - sts.chi2.cdf(LR_val, 5)
         print('chi squared of H0 with 5 degrees of freedom p-value:', pval_h0)
```

hypothesis value log likelihood: -2253.70068804

MLE log likelihood: 876.865046498

likelihood ratio value: 6261.13146908

chi squared of H0 with 5 degrees of freedom p-value: 0.0